

```
In [75]: !pip install seaborn
!pip install tensorflow --user
!pip install keras
!pip install daytime
!pip install torch
```

```
Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: seaborn in /home/student/.local/lib/
python3.10/site-packages (0.13.2)
Requirement already satisfied: pandas>=1.2 in /home/student/.local/
lib/python3.10/site-packages (from seaborn) (2.0.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /home/stu
dent/.local/lib/python3.10/site-packages (from seaborn) (3.7.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /home/studen
t/.local/lib/python3.10/site-packages (from seaborn) (1.24.3)
Requirement already satisfied: pillow>=6.2.0 in /usr/lib/python3/di
st-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (9.0.1)
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in /usr/lib/py
thon3/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.4.7)
Requirement already satisfied: kiwisolver>=1.0.1 in /home/studen
t/.local/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=
3.4->seaborn) (1.4.4)
Requirement already satisfied: cycycler>=0.10 in /home/student/.loca
l/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seabo
rn) (0.11.0)
```

```
In [76]: import pandas as pd
import numpy as np
import tensorflow as tf
```

```
In [77]: dataset = pd.read_csv("creditcard.csv")
```

```
In [78]: #dataset.head
print(dataset.columns)
dataset.describe()
```

```
Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9',
      'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19',
      'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
      'Class'],
      dtype='object')
```

Out[78]:

	Time	V1	V2	V3	V4	V5
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01

8 rows × 31 columns

In []:

```
In [79]: #checking for null value  
dataset.isnull().sum()
```

```
Out[79]: Time          0  
V1              0  
V2              0  
V3              0  
V4              0  
V5              0  
V6              0  
V7              0  
V8              0  
V9              0  
V10             0  
V11             0  
V12             0  
V13             0  
V14             0  
V15             0  
V16             0  
V17             0  
V18             0  
V19             0  
V20             0  
V21             0  
V22             0  
V23             0  
V24             0  
V25             0  
V26             0  
V27             0  
V28             0  
Amount          0  
Class           0  
dtype: int64
```


```
In [80]: print("No. of unique labels ", len(dataset['Class'].unique()))  
print("Label values ", dataset.Class.unique())  
  
No. of unique labels  2  
Label values  [0 1]
```

```
In [81]: #0 is for normal credit card transaction  
#1 is for fraudulent credit card transaction
```

```
In [82]: import matplotlib.pyplot as plt

# Count occurrences of each class
count_classes = dataset['Class'].value_counts()


# Plotting
count_classes.plot(kind='bar')
plt.title("Class Distribution")
plt.xlabel("Class")
plt.ylabel("Number of Observations")
plt.xticks(rotation=0)
plt.show()
```



```
In [83]: # Save the normal and fraudulent transactions in separate dataframe

normal_dataset = dataset[dataset.Class == 0]
fraud_dataset = dataset[dataset.Class == 1]
```

```
In [84]: #Visualize transaction amounts for normal and fraudulent transactions
bins = np.linspace(200, 2500, 100)
plt.hist(normal_dataset.Amount, bins=bins, alpha=1, density=True, label='Normal')
plt.hist(fraud_dataset.Amount, bins=bins, alpha=0.5, density=True, label='Fraud')
plt.legend(loc='upper right')
plt.title("Transaction amount vs Percentage of transactions")
plt.xlabel("Transaction amount (USD)")
plt.ylabel("Percentage of transactions");
plt.show()
```



```
In [85]: #Time and Amount are the columns that are not scaled, so applying StandardScaler
#Normalizing the values between 0 and 1 did not work great for the data
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
dataset['Time'] = sc.fit_transform(dataset['Time'].values.reshape(-1,1))
dataset['Amount'] = sc.fit_transform(dataset['Amount'].values.reshape(-1,1))
```

```
In [86]: #The last column in the dataset is our target variable.

from sklearn.model_selection import train_test_split
raw_data = dataset.values

# The last element contains if the transaction is normal which is represented by labels
labels = raw_data[:, -1]

# The other data points are the electrocardiogram data
data = raw_data[:, 0:-1]

train_data, test_data, train_labels, test_labels = train_test_split(data, labels,
```

```
In [ ]:
```

```
In [87]: #normalizing the training and testing data to scale the values between

min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)
train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)
train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)
```

```
In [88]: #Use only normal transactions to train the Autoencoder.
#Normal data has a value of 0 in the target variable. Using the target
train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

normal_train_data = train_data[~train_labels]
normal_test_data = test_data[~test_labels]
fraud_train_data = train_data[train_labels]
fraud_test_data = test_data[test_labels]
print(" No. of records in Fraud Train Data=",len(fraud_train_data))
print(" No. of records in Normal Train data=",len(normal_train_data))
print(" No. of records in Fraud Test Data=",len(fraud_test_data))
print(" No. of records in Normal Test data=",len(normal_test_data))
```

```
No. of records in Fraud Train Data= 402
No. of records in Normal Train data= 227443
No. of records in Fraud Test Data= 90
No. of records in Normal Test data= 56872
```

```
In [89]: nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1] #num of columns, 30
encoding_dim = 14
hidden_dim_1 = int(encoding_dim / 2) #
hidden_dim_2=4
learning_rate = 1e-7
```

```
In [90]: #input Layer
input_layer = tf.keras.layers.Input(shape=(input_dim, ))
```

```
In [91]: #Encoder
encoder = tf.keras.layers.Dense(encoding_dim, activation="tanh",activ
encoder=tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim_1, activation='relu')(enco
encoder = tf.keras.layers.Dense(hidden_dim_2, activation=tf.nn.leaky_
```

```
In [92]: # Decoder
decoder = tf.keras.layers.Dense(hidden_dim_1, activation='relu')(enco
decoder=tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim, activation='relu')(deco
decoder = tf.keras.layers.Dense(input_dim, activation='tanh')(decoder)
```

```
In [93]: #Autoencoder
autoencoder = tf.keras.Model(inputs=input_layer, outputs=decoder)
autoencoder.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 30)]	0
dense_6 (Dense)	(None, 14)	434
dropout_2 (Dropout)	(None, 14)	0
dense_7 (Dense)	(None, 7)	105
dense_8 (Dense)	(None, 4)	32
dense_9 (Dense)	(None, 7)	35
dropout_3 (Dropout)	(None, 7)	0
dense_10 (Dense)	(None, 14)	112
dense_11 (Dense)	(None, 30)	450

=====

Total params: 1168 (4.56 KB)
Trainable params: 1168 (4.56 KB)
Non-trainable params: 0 (0.00 Byte)

=====

```
In [94]: #Define the callbacks for checkpoints and early stopping
cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.h5")
```

```
In [95]: # define our early stopping
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=1,
    mode='min',
    restore_best_weights=True)
```

```
In [96]: #Compile the Autoencoder

autoencoder.compile(
    metrics=['accuracy'],
    loss='mean_squared_error',
    optimizer='adam')
```

```
In [97]: #Train the Autoencoder
history = autoencoder.fit(normal_train_data, normal_train_data,
epochs=nb_epoch,
batch_size=batch_size,
shuffle=True,

validation_data=(test_data, test_data),
verbose=1,

callbacks=[cp, early_stop]
).history

Epoch 1/50
3513/3554 [=====>.] - ETA: 0s - loss: 0.0042
- accuracy: 0.0325
Epoch 1: val_loss improved from inf to 0.00002, saving model to aut
oencoder_fraud.h5
3554/3554 [=====] - 4s 865us/step - loss:
0.0042 - accuracy: 0.0326 - val_loss: 2.0033e-05 - val_accuracy:
0.1247
Epoch 2/50
213/3554 [>.....] - ETA: 2s - loss: 1.8245
e-05 - accuracy: 0.0352

/home/student/.local/lib/python3.10/site-packages/keras/src/engine/
training.py:3000: UserWarning: You are saving your model as an HDF5
file via `model.save()`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g. `model.save
('my_model.keras')`.
  saving_api.save_model(
```



```
3483/3554 [=====>.] - ETA: 0s - loss: 1.9319
e-05 - accuracy: 0.0616
Epoch 2: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 3s 833us/step - loss:
1.9307e-05 - accuracy: 0.0618 - val_loss: 1.9989e-05 - val_accuracy
: 0.0110
Epoch 3/50
3486/3554 [=====>.] - ETA: 0s - loss: 1.9314
e-05 - accuracy: 0.0658
Epoch 3: val_loss did not improve from 0.00002
3554/3554 [=====] - 3s 828us/step - loss:
1.9352e-05 - accuracy: 0.0658 - val_loss: 2.0114e-05 - val_accuracy
: 0.2171
Epoch 4/50
3490/3554 [=====>.] - ETA: 0s - loss: 1.9375
e-05 - accuracy: 0.0692
Epoch 4: val_loss did not improve from 0.00002
3554/3554 [=====] - 3s 825us/step - loss:
1.9409e-05 - accuracy: 0.0689 - val_loss: 2.0084e-05 - val_accuracy
: 0.0418
Epoch 5/50
3514/3554 [=====>.] - ETA: 0s - loss: 1.8848
e-05 - accuracy: 0.1712
Epoch 5: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 3s 840us/step - loss:
1.8851e-05 - accuracy: 0.1723 - val_loss: 1.8681e-05 - val_accuracy
: 0.3232
Epoch 6/50
3513/3554 [=====>.] - ETA: 0s - loss: 1.8233
e-05 - accuracy: 0.2195
Epoch 6: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 3s 825us/step - loss:
1.8233e-05 - accuracy: 0.2187 - val_loss: 1.8656e-05 - val_accuracy
: 0.2306
Epoch 7/50
3550/3554 [=====>.] - ETA: 0s - loss: 1.8183
e-05 - accuracy: 0.1444
Epoch 7: val_loss did not improve from 0.00002
3554/3554 [=====] - 3s 832us/step - loss:
1.8182e-05 - accuracy: 0.1445 - val_loss: 1.9488e-05 - val_accuracy
: 0.2188
Epoch 8/50
3525/3554 [=====>.] - ETA: 0s - loss: 1.8202
e-05 - accuracy: 0.1272
Epoch 8: val_loss did not improve from 0.00002
3554/3554 [=====] - 3s 819us/step - loss:
1.8193e-05 - accuracy: 0.1276 - val_loss: 1.8851e-05 - val_accuracy
: 0.2266
Epoch 9/50
3521/3554 [=====>.] - ETA: 0s - loss: 1.7961
e-05 - accuracy: 0.1892
Epoch 9: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 3s 826us/step - loss:
1.7952e-05 - accuracy: 0.1898 - val_loss: 1.7849e-05 - val_accuracy
: 0.2472
Epoch 10/50
```

```


3519/3554 [=====>.] - ETA: 0s - loss: 1.7411
e-05 - accuracy: 0.2809
Epoch 10: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 3s 823us/step - loss:
1.7418e-05 - accuracy: 0.2811 - val_loss: 1.7601e-05 - val_accuracy
: 0.3152
Epoch 11/50
3505/3554 [=====>.] - ETA: 0s - loss: 1.7296
e-05 - accuracy: 0.2915
Epoch 11: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
Restoring model weights from the end of the best epoch: 1.
3554/3554 [=====] - 3s 829us/step - loss:
1.7306e-05 - accuracy: 0.2916 - val_loss: 1.7454e-05 - val_accuracy
: 0.3497
Epoch 11: early stopping

```

```

In [98]: #Plot training and test loss
plt.plot(history['loss'], linewidth=2, label='Train')
plt.plot(history['val_loss'], linewidth=2, label='Test')
plt.legend(loc='upper right')
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
#plt.ylim(ymin=0.70,ymax=1)
plt.show()

```



```


In [99]: #Detect Anomalies on test data
'''
Anomalies are data points where the reconstruction loss is higher
To calculate the reconstruction loss on test data, predict the test
data and calculate the mean square error between the test data and th
'''
test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2), axis=1)
error_df = pd.DataFrame({'Reconstruction_error': mse, 'True_class': te

1781/1781 [=====] - 1s 458us/step

```

```
In [100]: #Plotting the test data points and their respective reconstruction error
#if the threshold value needs to be adjusted.
threshold_fixed = 0.001
groups = error_df.groupby('True_class')
fig, ax = plt.subplots()
for name, group in groups:

    ax.plot(group.index, group.Reconstruction_error, marker='o', ms=3)
ax.hlines(threshold_fixed, ax.get_xlim()[0], ax.get_xlim()[1], color='red')
ax.legend()
plt.title("Reconstruction error for normal and fraud data")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show();
```



```
In [106]: '''Detect anomalies as points where the reconstruction loss is greater than
a fixed threshold.
Here we see that a value of 52 for the threshold will be good.
Evaluating the performance of the anomaly detection'''

from sklearn.metrics import confusion_matrix, recall_score, accuracy_score
import seaborn as sns


threshold_fixed = 0.001
pred_y = [1 if e > threshold_fixed else 0 for e in error_df.Reconstruction_loss]
error_df['pred'] = pred_y

conf_matrix = confusion_matrix(error_df.True_class, pred_y)

plt.figure(figsize=(4, 4))

LABELS = ["Normal", "Fraud"]

sns.heatmap(conf_matrix, xticklabels=LABELS, yticklabels=LABELS, annot=True)
plt.title("Confusion matrix")
plt.ylabel('True class')
plt.xlabel('Predicted class')
plt.show()
# print Accuracy, precision and recall
print(" Accuracy: ", accuracy_score(error_df['True_class'], error_df['pred']))
print(" Recall: ", recall_score(error_df['True_class'], error_df['pred']))
print(" Precision: ", precision_score(error_df['True_class'], error_df['pred']))
```



```
Accuracy: 0.9980688880306169
Recall: 0.15555555555555556
Precision: 0.2916666666666667
```

In []: